

Patterns for Cleaning Up Bug Data

Rodrigo Souza^{*†}, Christina Chavez^{*‡}, and Roberto Bittencourt^{*§}

^{*}Dept. of Computer Science, Federal University of Bahia, Brazil

[†]Data Processing Center, Federal University of Bahia, Brazil

[‡]Fraunhofer Project Center, Federal University of Bahia, Brazil

[§]Dept. of Exact Sciences, State University of Feira de Santana, Brazil

rodrigo@dcc.ufba.br, flach@dcc.ufba.br, roberto@uefs.br

Abstract—Bug reports provide insight about the quality of an evolving software and about its development process. Such data, however, is often incomplete and inaccurate, and thus should be cleaned before analysis. In this paper, we present patterns that help both novice and experienced data scientists to discard invalid bug data that could lead to wrong conclusions.

Index Terms—data analysis, mining software repositories, patterns, bugs.

I. INTRODUCTION

Bug tracking systems store bug reports for a software project and keep record of discussion and progress changes for each bug. This data can be used not only to assess quality attributes of the software, but also to reason about its development process. Such richness of information makes bug reports an invaluable source for data scientists interested in software engineering.

However, bug tracking systems often contain data that is inaccurate, incomplete [1], or biased [2]. For example, changing the status of a bug report to `VERIFIED` usually means that, after a resolution was found to the bug, some kind of software verification (e.g., source code inspection, testing) was performed and the resolution was considered appropriate. Sometimes, however, old bug reports are marked as verified just to help users and developers keep track of current bug reports [3].

Without proper guidance, it is easy to overlook pitfalls in the data and draw wrong conclusions. In this paper, we provide best practices and step-by-step solutions to recurring problems related to cleaning bug data.

II. DATA SET

Each pattern contains an *Examples* section with code snippets showing how to apply the pattern on real data. The snippets are written in R, a programming language for data analysis¹. The data used are bug reports from the project NetBeans/Platform, made available for the 2011 edition of the MSR Mining Challenge².

The NetBeans project uses Bugzilla³ as its bug tracking system, which stores all data in a MySQL database. The source code presented here refers to database tables and columns used

TABLE I
SAMPLE OF `changes` TABLE, HOUR INFO OMITTED.

bug	user	time	field	new.value
427	17822	2009-10-30	resolution	WONTFIX
500	182	2002-04-12	bug_status	CLOSED
755	182	2002-01-11	bug_status	REOPENED

by Bugzilla, but it should work with any bug tracking system with minor changes.

Although the full NetBeans data set contains 57 database tables, in this paper only three are used: `bugs`, `changes` (originally, `bugs_activity`), and `comments` (originally, `longdescs`).

The `bugs` table contains general information about each bug report, which is identified by a unique number (column `bug`). Each bug report has a severity, a priority, and two timestamps: the time of creation (`creation.time`), and the time of the last modification (`modif.time`).

The `comments` table contains comments that each user added to a bug report at some point in time. To reduce the file size, the comment text was replaced by its MD5 hash (column `comment.md5`). With high probability, two comments are represented by the same hash number if and only if they contain the same text.

The `changes` table contains all modifications users made on bug reports over time. This includes changes in priority, status, resolution, or any other field in a bug report. Each row contains the new value of a field that was modified by a user⁴ at some point in time. Table I shows a sample of the `changes` table.

The `bug_status` field is used to track the progress of the bug fixing activity. A bug report is created with status `NEW` or `UNCONFIRMED`. Then, its status may be changed to `ASSIGNED`, to denote that a user has taken responsibility on the bug. After that, the bug is `RESOLVED`, then optionally `VERIFIED` by the quality assurance team, and finally `CLOSED` when the next software release comes out. If, after resolving the bug, someone finds that the resolution was not appropriate, the status is changed to `REOPENED`.

All the data and code used in this paper is online⁵.

¹<http://www.r-project.org/>

²<http://2011.msrconf.org/msr-challenge.html>

³<http://www.bugzilla.org/>

⁴In this context, user denotes a user of the bug tracking system, which can be either a developer or a final user.

⁵<https://github.com/rodrigors/dapse13-bugpatterns>

III. LOOK OUT FOR MASS UPDATES

A. Problem

Determine which changes to bug reports were the result of a mass update.

B. Context

Changes to a bug report are often the result of an effort made by developers to triage, fix or verify a bug. Sometimes, however, hundreds or thousands of bug reports are changed almost simultaneously. Such changes are not caused by a burst of productivity; instead, they are the result of a mass update, often performed with the purpose of cleaning up the bug tracking system.

Mass updates can also be motivated by a policy change. In Eclipse Modeling Framework, for instance, developers decided that bug reports containing fixes that were already published on their website should have the status `VERIFIED`⁶. A mass status update was needed to make previous bug reports conform to the new policy.

Mass updates are characterized by a large number of changes of the *same type* (e.g., marking a bug report as `VERIFIED` or changing a target milestone) made by a single developer in a short period of time. Often such changes are accompanied by a comment that is the same for all changes in the mass update.

C. Solution

First, choose the type of change that you wish to analyze (e.g., changing a bug status to `VERIFIED`). Then apply one of the following solutions.

Solution 1. Select only the changes of the chosen type, along with the time of the change. Sort the changes by time and then plot the accumulated number of changes over time as a line chart.

The line is monotonically increasing, but periods with a large number of changes will stand out as steep ascents. Examine the vertical axis to assess whether such ascents represent a large number of changes (e.g., thousands). If this is the case, then it is likely that the changes were caused by a mass update.

Solution 2. Select only the changes of the chosen type, along with the date of the change, the user who made the change and the comment. Then, group by similar triples $\langle \text{date}, \text{user}, \text{comment} \rangle$, and sort them by number of occurrences. Triples with frequency above a given threshold (e.g., hundreds) are good candidates for mass updates. You may read the comment text and look for references to cleanup, policy change or mass update.

D. Discussion

The first solution is useful for a quick and exploratory assessment of mass updates, but it is subjective and may be inaccurate. For example, some projects have testing phases, in which the verification activity is higher than usual. Such




Fig. 1. Accumulated number of changes over time.

phases can be identified by applying the *Testing Phase* pattern [4]. In this case, the bug status is not mass updated to `VERIFIED`; instead, each change reflects an actual verification effort.

In the second solution, users and their comments are considered, which helps improve its accuracy. It is a numeric, objective solution, but it is also more computationally intensive.

E. Examples

Souza and Chavez [3] used the first solution (see Figure 2 in their paper) and a variation of the second without taking comments into account to detect mass verifications (i.e., changes which set the status to `VERIFIED`). They discarded all changes that were part of mass updates which changed at least 50 bug reports.

Solution 1. Here is a sample R code to plot the accumulated number of verifications over time, which produces Figure 1.

```
> ver <- subset(changes, field == "bug_status"
+   & new.value == "VERIFIED")
> ver <- ver[order(ver$time), ]
> ver$n.changes <- 1:nrow(ver)
> with(ver, plot(n.changes ~ time, type="l"))
```

In this chart, some line segments are almost vertical (e.g., the line between 2002 and 2004). Such segments mark dates with mass updates.

Solution 2. Here is a sample R code to count updates by user, date, and comment. First, select only verifications, and then use the `merge` operation to associate them with their respective comments. After that, group and count them as usual. The 6 records with the highest counts are shown in Table II.

```
> library(plyr)
> ver <- subset(changes, field == "bug_status"
+   & new.value == "VERIFIED")
> ver$date <- as.Date(ver$time)
> full <- merge(ver, comments)
> cnt <- count(full, c("date", "user",
+   "comment.md5"))
> cnt <- cnt[order(cnt$freq, decreasing=T), ]
```

⁶See http://wiki.eclipse.org/Modeling_PMC_Meeting_2007-10-16

TABLE II
SAMPLE OF MASS UPDATES, COMMENTS OMITTED.

date	user	freq
2003-07-01	17822	1703
2003-07-01	17822	972
2003-07-02	17822	447
2005-07-12	182	437
2005-12-20	182	209
2005-07-13	182	181

IV. OLD WINE TASTES BETTER

A. Problem

Determine bug reports that are too recent to be classified.

B. Context

Bug reports change over time. Sometimes, one needs to classify a bug report according to the eventual occurrence of some change. For example, suppose that one wants to predict whether future bug reports will be reopened. To train a prediction model, each existing bug report has to be classified as reopening (if it was or will be reopened) or non-reopening (if it will likely never be reopened). However, for recent bug reports, it is likely that they were not reopened yet—even if they will be, eventually. Therefore, recent reports cannot be accurately classified with available data, and should be discarded from the training set. The question is how long one should wait for a change to happen before assuming it will never happen.

C. Solution

Assume you want to classify each bug report, according to a particular change (e.g., reopening, fixing), as *eventually* (if the change happened or will happen, eventually) or *never* (if the change will likely never happen). If the bug report has already undergone the change, classify it as *eventually*. If not, then measure its lifetime, from creation to the last date available in the data. If the lifetime is long enough (i.e., it is above some *threshold*), then classify it as *never*. If the lifetime is short, however, one cannot be confident that the change will never happen, hence discard it from the analysis.

The key is to choose an appropriate threshold for the lifetime of bug reports. At first, just choose any value (e.g., 30 days). Then, compute the confidence, α , that bugs older than the threshold will never undergo the change. To do that, first compute the proportion, c of bug reports that have undergone the change. Then, compute the proportion of older bug reports (i.e., which have a lifetime greater than the threshold), t , that did not undergo the change within the threshold. The confidence, α , can be approximated by the quotient c/t .

Assess whether α is high enough (most data scientists find 0.95 to be an acceptable value). If it is not, choose another

threshold and recompute α until you are satisfied. After finding an appropriate threshold, discard bugs with lifetime shorter than the threshold, because it is not possible to determine, with confidence α , that they will never undergo the change.

D. Discussion

It is a common mistake to keep recent bug reports. This is equivalent to choosing a threshold of 0, a value that is too optimistic. There is a trade-off between the confidence, α , and the size of the final data set: the higher the confidence, the higher the number of discarded bugs.

E. Examples

Here's how to apply this pattern using R to analyze the change of a bug status to REOPENED. First of all, create a data frame `data`, augmenting bugs with information about their first reopening and their lifetime.

```
> library(data.table)
> reopenings <- data.table(changes)[
+   field == "bug_status" & new.value == "REOPENED",
+   list(time.first.reopen = min(time)), by=bug]
> data <- merge(bugs, reopenings, all.x=T)
> data$days.to.reopen <- as.numeric(
+   data$time.first.reopen - data$creation.time,
+   units="days")
> last.time <- max(data$modif.time)
> data$lifetime <- as.numeric(
+   last.time - data$creation.time,
+   units="days")
```

In this example, we use a threshold of 42 days (6 weeks). Use it to compute α (alpha):

```
> threshold <- 42
> older <- subset(data, lifetime > threshold)
> c <- sum(is.na(data$days.to.reopen)) / nrow(data)
> t <- sum(is.na(older$days.to.reopen)
+   | older$days.to.reopen > threshold) / nrow(older)
> alpha <- c / t
```

In this case, we find $\alpha = 0.95$, an acceptable value. Therefore, you should only analyze bug reports older than 42 days—which are already stored in the variable `older`—and discard the rest. In this case, only 0.7% of the bugs needed to be discarded.

REFERENCES

- [1] J. Aranda and G. Venolia, "The secret life of bugs: Going past the errors and omissions in software repositories," in *Proc. of the 31st Int. Conf. on Soft. Engineering*, 2009, pp. 298–308.
- [2] C. Bird, A. Bachmann, E. Aune, J. Duffy, A. Bernstein, V. Filkov, and P. Devanbu, "Fair and balanced?: bias in bug-fix datasets," in *European Soft. Eng. Conf. and Symposium on the Foundations of Soft. Eng.*, ser. ESEC/FSE '09. ACM, 2009.
- [3] R. Souza and C. Chavez, "Characterizing verification of bug fixes in two open source IDEs," in *Proceedings of the 9th Working Conference on Mining Software Repositories*. IEEE, June 2012.
- [4] R. Souza, C. Chavez, and R. Bittencourt, "Patterns for extracting high level information from bug reports," in *Proceedings of the First Workshop on Data Analysis Patterns in Software Engineering*. IEEE, May 2013.