Patterns for cleaning bug data

Rodrigo Souza*†, Christina Chavez*‡ and Roberto Bittencourt*§
*Dept. of Computer Science, Federal University of Bahia, Brazil
†Data Processing Center, Federal University of Bahia, Brazil
‡Fraunhofer Project Center, Federal University of Bahia, Brazil
§Dept. of Exact Sciences, State University of Feira de Santana, Brazil
Email: rodrigo@dcc.ufba.br, flach@dcc.ufba.br, roberto@uefs.br

Abstract—**Bug reports provide insight about the quality of an evolving software and about its development process. Such data, however, is often incomplete and inaccurate, and thus should be cleaned before analysis. In this paper, we present patterns that help both novice and experienced data scientists to discard invalid bug data that could lead to wrong conclusions.**

Index Terms—data analysis, mining software repositories, patterns, bugs.

I. Introduction

Bug tracking systems store bug reports for a software project and keep record of discussion and progress changes for each bug. This data can be used not only to assess quality attributes of the software, but also to reason about its development process. Such richness of information makes bug reports an invaluable source for data scientists interested in software engineering.

However, bug tracking systems often contain data that is inaccurate [1], biased [2], or incomplete [3]. For example, changing the status of a bug report to VERIFIED usually means that, after a resolution was found to the bug, some kind of software verification (source code inspection, testing etc.) was performed and the resolution was considered appropriate. Sometimes, however, old bug reports are marked as verified just to help users and developers keep track of current bug reports [4].

Without proper guidance, it is easy to overlook pitfalls in the data and draw wrong conclusions. In this paper, we provide best practices and step-by-step solutions to recurring problems related to cleaning bug data.

Each solution is presented in a structured form called pattern. A pattern consists of 6 sections: (1) a short *name*; (2) the *problem* being solved; (3) a *context* in which the pattern can be applied; (4) one or more *solutions* to the problem; (5) a *discussion* of trade-offs and common mistakes to consider when using the pattern; and (6) *examples* of the pattern in use.

II. Data Set

Each pattern contains an *Example* section with code snippets showing how to apply the pattern on real data. The snippets are written in R, a programming language for data analysis[1]. The data used are the bug reports from the project NetBeans/Platform, made available for the 2011 edition of the Mining Challenge[2].

The NetBeans project uses Bugzilla[3] as its bug tracking system, and stores all data in a MySQL database. Although the source code presented here refers to database tables and columns used by Bugzilla, it should work with any bug tracking system with minor modifications.

Although the full NetBeans data set contains 57 database tables, in this paper only three are used: `bugs`, `changes` (originally `bugs_activity`), and `comments` (originally `longdescs`). Notice that some tables and columns were renamed for clarity purposes.

The `bugs` table contains general information about each bug report, which is identified by a unique number (column `bug`). Each bug report has a `severity`, a `priority`, and two timestamps: the time of creation (`creation.time`), and the time of the last modification (`modif.time`). Table I shows a sample of the `bugs` table.

TABLE I
Sample of **bugs** table, omitting the hour part of the times.

bug	severity	priority	creation.time	modif.time
397	normal	P4	1998-07-31	2008-12-23
427	normal	P4	1998-08-08	2008-12-23
479	normal	P2	1998-08-19	2008-12-23

The `comments` table contains comments that each `user` added to a `bug` report at some point in `time`. To reduce the file size, the text for the comment was replaced by its MD5 hash (column `comment.md5`). With high probability, two comments are represented by the same hash number if and only if they contain the same text. Table II shows a sample of the `comments` table.

TABLE II
Sample of **comments** table, omitting the hour part of **time**.

bug	comment.md5	user	time
397	bafe80ea4ba2b73b6883243d8718ae3b	1887	2000-11-01
427	64d38f33f8848525a81b415406d064df	46	1998-08-08
427	264e61d606f0a4c0c8d477de00fba347	124	2000-07-25

---

[1]http://www.r-project.org/

[2]http://2011.msrconf.org/msr-challenge.html
[3]http://www.bugzilla.org/

The `changes` table contains all modifications made by users on bug reports over time. This includes changes in priority, status, resolution, and virtually any other field in a bug report. Each row contains the `new value` of a `field` that was modified by a `user`[4] at some point in `time`. Table III shows a sample of the `bugs` table.

TABLE III
Sample of `changes` table, omitting the hour part of `time`.

bug	user	time	field	new.value
427	17822	2009-10-30	resolution	WONTFIX
500	182	2002-04-12	bug_status	CLOSED
755	182	2002-01-11	bug_status	REOPENED

The *status* field is used to keep track of the progress of the bug fixing activity. A bug report is created with status `NEW` or `UNCONFIRMED`. Then, its status may be changed to `ASSIGNED`, to denote that a user has taken responsibility on the bug. After that, the bug is `RESOLVED`, then optionally `VERIFIED` by the quality assurance team and `CLOSED`, after the next software release comes out. If, after resolving the bug, someone finds that the resolution was not appropriate, the status is changed to `REOPENED`.

There are many forms of resolving a bug. To reflect that, when a bug status is changed to `RESOLVED`, the *resolution* field is changed either to `FIXED`—if the software was changed to solve the issue—, `WORKSFORME`—if developers were not able to reproduce the problem—, `DUPLICATE`—if a previous bug report describes the same problem—, among other resolutions.

All the data and code used in this paper is available online[5].

III. Look Out For Mass Updates

Problem
Discover changes to bug reports that were the result of a mass update.

Context
Changes to a bug report are often the result of an effort made by developers to triage, fix or verify a bug. Sometimes, however, hundreds or thousands of bug reports are changed almost simultaneously. Such changes are not caused by a burst of productivity; instead, they are the result of a mass update, often performed with the purpose of cleaning up the bug tracking system.

Mass updates can also be motivated by a policy change. In Eclipse Modeling Framework, for instance, developers decided that bug reports containing fixes that were already published on their website should have the status `VERI-FIED`[6]. A mass status update was needed to make previous bug reports conform to the new policy.

[4]In this context, user denotes a user of the bug tracking system, which can be either a developer or a final user.
[5]https://github.com/rodrigorgs/dapse13-bugpatterns
[6]Source: http://wiki.eclipse.org/Modeling_PMC_Meeting,_2007-10-16

Mass updates are characterized by a large number of changes of the *same type* (e.g., marking a bug report as `VERIFIED` or changing a target milestone) made by a single developer in a short period of time. Often such changes are accompanied by a comment that is the same for all changes in the mass update.

Solution

First, choose the type of change that you wish to analyze (e.g., changing a bug status to `VERIFIED`). Then apply one of the following solutions.

Solution 1. Select only the changes of the chosen type, along with the time of the change. Sort the changes by time and then plot the accumulated number of changes over time as a line chart.

The line is monotonically increasing, but periods with a large number of changes will stand out as steep ascents. Examine the vertical axis to assess whether such ascents represent a large number of changes (e.g., thousands). If this is the case, then it is likely that the changes were caused by a mass update.

Solution 2. Select only the changes of the chosen type, along with the date of the change, the user who made the change and the comment. Then, group by similar triples ⟨*date, user, comment*⟩, and sort them by number of occurrences. Triples with frequency above a given threshold (e.g., hundreds) are good candidates for mass updates. You may read the comment text and look for references to cleanup, policy change or mass update.

Discussion

The first solution to find mass updates is more visual, but less accurate, as it does not take into account the user who made the changes. It is useful for a quick exploratory assessment of mass updates. The second solution is numeric and takes into account the users and their comments, but is more computationally intensive.

Examples

For example, Souza and Chavez [4] used the first solution (see Figure 2 in their paper) and a variation of the second without taking comments into account to detect mass verifications (i.e., changes whose status is set to `VERIFIED`). They discarded all changes that were part of a mass update which updated at least 50 bug reports.

Solution 1. Here is a sample R code to plot the accumulated number of verifications over time.

```
> ver <- subset(changes, field == "bug_status"
+    & new.value == "VERIFIED")
> ver <- ver[order(ver$time), ]
> ver$n.changes <- 1:nrow(ver)
> with(ver, plot(n.changes ~ time, type="l"))
```

In this chart, some line segments are almost vertical (e.g., the line between 2002 and 2004). Such segments mark dates with mass updates.

Solution 2. Here is a sample R code to count updates by user, date, and comment. First, select only verifications, and then use the `merge` operation to associate them with their respective comments. After that, group and count them as usual. The 6 records with the highest counts are shown in Table IV.

```
> library(plyr)
> ver <- subset(changes, field == "bug_status"
+   & new.value == "VERIFIED")
> ver$date <- as.Date(ver$time)
> full <- merge(ver, comments)
> cnt <- count(full, c("date", "user",
+                       "comment.md5"))
> cnt <- cnt[order(cnt$freq, decreasing=T), ]
> head(cnt[, c("date", "user", "freq")])
```

TABLE IV
Sample of mass updates, comments ommitted.

date	user	freq
2003-07-01	17822	1703
2003-07-01	17822	972
2003-07-02	17822	447
2005-07-12	182	437
2005-12-20	182	209
2005-07-13	182	181

IV. Old Wine Tastes Better

Problem

Discard bug reports that are too recent to be classified.

Context

Bug reports change over time. Sometimes, one needs to classify a bug report according to the eventual occurence of some change. For example, has a bug been fixed (or will it, eventually)? Is this a bug that was or will be reopened, or it will never be reopened?

While for older bugs such questions can be answered with high confidence, for bugs that are too recent, the change we are interested may not have happened—*yet*. The question is how long to wait for the change to happen before assuming that it will never happen for a particular bug report.

Solution

For the sake of example, consider we are studying bug reopenings. If the data shows that a bug report was reopened some day, classify it as a reopening bug. If not, then measure its lifetime, from creation to the last date available in the data. If the lifetime is long enough (i.e., it is above some *threshold*), then classify it as non-reopening. That means the bug was not and, with high confidence, will never be reopened. If the lifetime is short, however, it is hard to predict whether the bug report will eventually reopen, hence discard it from the analysis.

So, first of all, choose a threshold. Then, compute the probability that, given that a bug report did not reopen within the threshold, it will never reopen (let's call it α, or confidence). To do that, compute the proportion of bug reports that were reopened, r. Then, compute the proportion of older bug reports (i.e., which have a lifetime greater than the threshold) that were not reopened within the threshold, t. The probability α can be approximated by the quotient $\frac{r}{t}$.

Assess whether the probability α is high enough (most data scientists find 0.95 to be an acceptable value). If it is not, choose another threshold and recompute α until you are satisfied. After finding an appropriate threshold, discard bugs with lifetime shorter than the threshold.

Discussion

It is a common mistake to keep recent bug reports. This is equivalent to choosing a threshold of 0, a value that is too optimistic.

Examples

Here's how to apply this pattern using R to analyze bug reopening. First of all, create a data frame `data`, augmenting `bugs` with information about their first reopening and their lifetime.

```
> library(data.table)
> reopenings <- data.table(changes)[
+   field == "bug_status" & new.value == "REOPENED",
+   list(time.first.reopen = min(time)), by=bug]
> data <- merge(bugs, reopenings, all.x=T)
> data$days.to.reopen <- as.numeric(
+   data$time.first.reopen - data$creation.time,
+   units="days")
> last.time <- max(data$modif.time)
> data$lifetime <- as.numeric(
+   last.time - data$creation.time,
+   units="days")
```

In this example, we use a threshold of 42 days (6 weeks). Use it to compute α (`alpha`):

```
> threshold <- 42
> older <- subset(data, lifetime > threshold)
> r <- sum(is.na(data$days.to.reopen)) / nrow(data)
> t <- sum(is.na(older$days.to.reopen)
+   | older$days.to.reopen > threshold) / nrow(older)
> alpha <- r / t
```

In this case, we find $\alpha = 0.95$, which is acceptable. Therefore, you should analyze only bug reports older than 42 days—which are already stored in the variable `older`—

and discard the rest. In this case, 0.7% of the bugs needed to be discarded.

References

[1] G. Antoniol, K. Ayari, M. Di Penta, F. Khomh, and Y.-G. Guéhéneuc, "Is it a bug or an enhancement?: a text-based approach to classify change requests," in *Proc. of the 2008 Conf. of the Center for Adv. Studies on Collaborative Research*, ser. CASCON '08.   ACM, 2008, pp. 23:304–23:318.

[2] C. Bird, A. Bachmann, E. Aune, J. Duffy, A. Bernstein, V. Filkov, and P. Devanbu, "Fair and balanced?: bias in bug-fix datasets," in *European Soft. Eng. Conf. and Symposium on the Foundations of Soft. Eng.*, ser. ESEC/FSE '09.   ACM, 2009.

[3] J. Aranda and G. Venolia, "The secret life of bugs: Going past the errors and omissions in software repositories," in *Proc. of the 31st Int. Conf. on Soft. Engineering*, 2009, pp. 298–308.

[4] R. Souza and C. Chavez, "Characterizing verification of bug fixes in two open source ides." in *MSR*, M. Lanza, M. D. Pent, and T. Xi, Eds.   IEEE, 2012, pp. 70–73. [Online]. Available: http://dblp.uni-trier.de/db/conf/msr/msr2012.html#SouzaC12